

# Nathan Teshome

Cambridge, Massachusetts | [nnt@mit.edu](mailto:nnt@mit.edu) | +1 617 719 0373 | [linkedin.com/in/nathan-teshome-84b675325](https://linkedin.com/in/nathan-teshome-84b675325) | [github.com/nnt-git13](https://github.com/nnt-git13)

## EDUCATION

### Massachusetts Institute of Technology (MIT)

B.S. Electrical Engineering; B.S. Mathematics; GPA: 4.60/5.00

Cambridge, MA

Anticipated 2028

- Pursuing NEET Autonomous Machines Certification.
- **Currently Enrolled:** Software Construction; Computation Structures; Semiconductor Electronic Circuits; Fundamentals of Statistics; C & Assembly.
- **Completed:** Algorithms I; Probability & Random Variables; Machine Learning; Circuits; Differential Equations; Linear Algebra; Electricity & Magnetism; Fundamentals of Programming (Python).

## ACADEMIC & PERSONAL PROJECTS

### FlashFlow: GPU-Optimized Transformer Inference Runtime

GitHub: [github.com/nnt-git13/flashflow](https://github.com/nnt-git13/flashflow)

- Trained a small GPT-style Transformer achieving <0.5 perplexity gap vs a strong PyTorch baseline on held-out data.
- Implemented fused CUDA/Triton kernels for attention and MLP using tiling and shared-memory blocking, yielding 1.4–1.9× kernel-level speedups in internal benchmarks.
- Built a C++ scheduling runtime lowering FX graphs into fused GPU operators with auto-tuned tile sizes; improved tokens/sec by 1.3–1.6× vs PyTorch eager.
- Integrated FP16/BF16 quantization and an early INT8 path; observed 2–3× higher arithmetic throughput with minimal accuracy degradation.

### RISC-V Vector CPU + Quant Finance Accelerator + Custom PCB

GitHub: [github.com/nnt-git13/riscv\\_quant.accel](https://github.com/nnt-git13/riscv_quant.accel)

- Built a functional RV32I soft core with preliminary custom vector instructions and a memory-mapped accelerator interface on an Artix-7 FPGA.
- Implemented a Monte-Carlo pricing accelerator with SIMD-style parallelism, fixed-point numerics, and reduction trees; achieved order-of-magnitude speedups vs Python/C++ baselines.
- Designed a 4-layer PCB including regulators, level-shifting, and PMOD expansion; validated core functionality through SystemVerilog benches and golden-model testing.

### Mini-TPU: MLIR Compiler + Programmable Tensor-Core Accelerator

GitHub: [github.com/nnt-git13/TODO](https://github.com/nnt-git13/TODO)

- Developing a programmable tensor-core accelerator (systolic array + DMA engine) with a custom MLIR dialect supporting tiling, fusion, and quantization-aware lowering.
- Early FPGA prototypes validate correctness with small FP16/BF16 workloads; initial tests show 2–3× throughput gains on matrix-multiply kernels.

## WORK EXPERIENCE

### MIT Brain-Score AI Neuroscience Benchmarking Platform

Undergraduate Software Engineer — Leaderboard Infrastructure

Cambridge, MA

Jun. 2025 – Ongoing

- Redesigned the Brain-Score model leaderboard (Python, Django, AG Grid) to support timestamp-based filtering, multi-index score aggregation, and cross-model comparison; improved load latency by ~35%.
- Implemented a dynamic Wayback Slider for historical model evaluation, wiring dual-date sliders to Score-level metadata and adding backend API routes for time-bounded queries.
- Built new aggregation pipelines for recursive benchmark hierarchies (vision & language), adding BFS tree construction, bottom-up score fusion, and robust handling of missing leaf scores.

### Satellite Swarm Reinforcement Learning Control — MIT

Undergraduate Researcher

Cambridge, MA

Nov. 2026 – Ongoing

- Ported EMFF swarm dynamics + RL controller into a GPU-parallel pipeline using vectorized integrators; increased simulation throughput by 8–12× on mid-range GPUs.
- Redesigned PPO/SAC components with custom rollout buffers and batched normalization, reducing training wall-time from multiple days to under 10 hours in internal tests.
- Prototyped attention-based multi-agent policies (transformer + GNN hybrids) achieving more stable convergence and reduced  $\Delta V$  usage across held-out orbital scenarios.

## TECHNICAL SKILLS

**Programming:** Python, C, C++, SystemVerilog/Verilog, Bash; PyTorch, NumPy/Pandas, LLVM/MLIR; ROS 2, OpenCV; Django/REST; JavaScript; CI/CD (pytest, CMake, GitHub Actions).

**Hardware:** FPGA (Zynq-7000/Artix-7/ZCU104), Vivado/Vitis, PYNQ; AXI4/AXI-Lite/AXI-Stream, BRAM/DSP48; cocotb, riscv-dv, SymbiYosys; oscilloscopes, logic analyzers, JTAG; KiCad PCB design, soldering.

**Tools:** Linux, Docker, Git, SLURM, RViz2, Onshape, 3D printing/rapid prototyping.